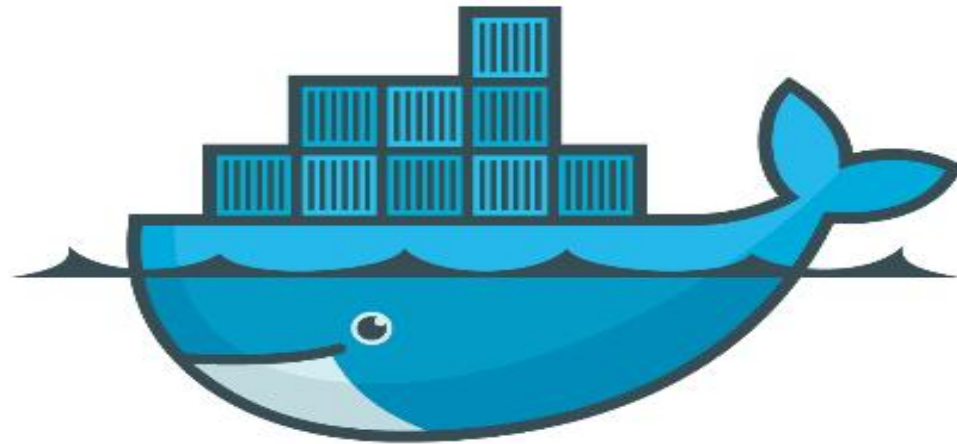


Case Study: OS Hot Topics

Instructor: Dr. Liting Hu

Introduction to Docker



docker

What is Docker?

- Docker is an open-source project that automates the deployment of applications inside software containers, by providing an additional layer of **abstraction** and **automation** of **operating system–level virtualization** on Linux.

Docker History

- A dotCloud (PAAS provider) project
- Initial commit January 18, 2013
- Docker 0.1.0 released March 25, 2013
- 18,600+ github stars, 3800+ forks, 740 Contributors.... and continues
- dotCloud pivots to docker inc. October 29, 2013

Docker Features

- Light-Weight
 - Minimal overhead (*cpu/io/network*)
 - Based on Linux containers
 - Uses layered filesystem to save space (AUFS/LVM)
 - Uses a copy-on-write filesystem to track changes
- Portable
 - Can run on any Linux system that supports LXC (today).
 - 0.7 release includes support for RedHat/Fedora family.
 - Raspberry pi support.
 - Future plans to support other container tools (lmctfy, etc.)
 - Possible future support for other operating systems (Solaris, OSX, Windows?)
- Self-sufficient
 - A Docker container contains everything it needs to run
 - Minimal Base OS
 - Libraries and frameworks
 - Application code
 - A docker container should be able to run anywhere that Docker can run.

The Challenge


Multiplicity of Stacks

 Static website
nginx 1.5 + modsecurity + openssl + bootstrap 2


 User DB
postgresql + pgv8 + v8

 Queue
Redis + redis-sentinel

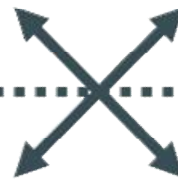
 Analytics DB
hadoop + hive + thrift + OpenJDK

 Background workers
Python 3.0 + celery + pyredis + libcurl + ffmpeg + libopencv + nodejs + phantomjs

 Web frontend
Ruby + Rails + sass + Unicorn

 API endpoint
Python 2.7 + Flask + pyredis + celery + pycopg + postgresql-client

Do services and apps
interact
appropriately?



Public Cloud

Production Cluster



Disaster recovery

Contributor's laptop



Production Servers

Can I migrate
smoothly and
quickly?

Development VM



QA server




Customer Data Center



The Matrix From Hell



Static website	?	?	?	?	?	?	?
Web frontend	?	?	?	?	?	?	?
Background workers	?	?	?	?	?	?	?
User DB	?	?	?	?	?	?	?
Analytics DB	?	?	?	?	?	?	?
Queue	?	?	?	?	?	?	?
	Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers



Cargo Transport Pre-1960

Multiplicity of Goods
















Do I worry about
how goods interact
(e.g. coffee beans
next to spices)

Multiplicity of
methods for
transporting/storing



Can I transport quickly
and smoothly
(e.g. from boat to train
to truck)

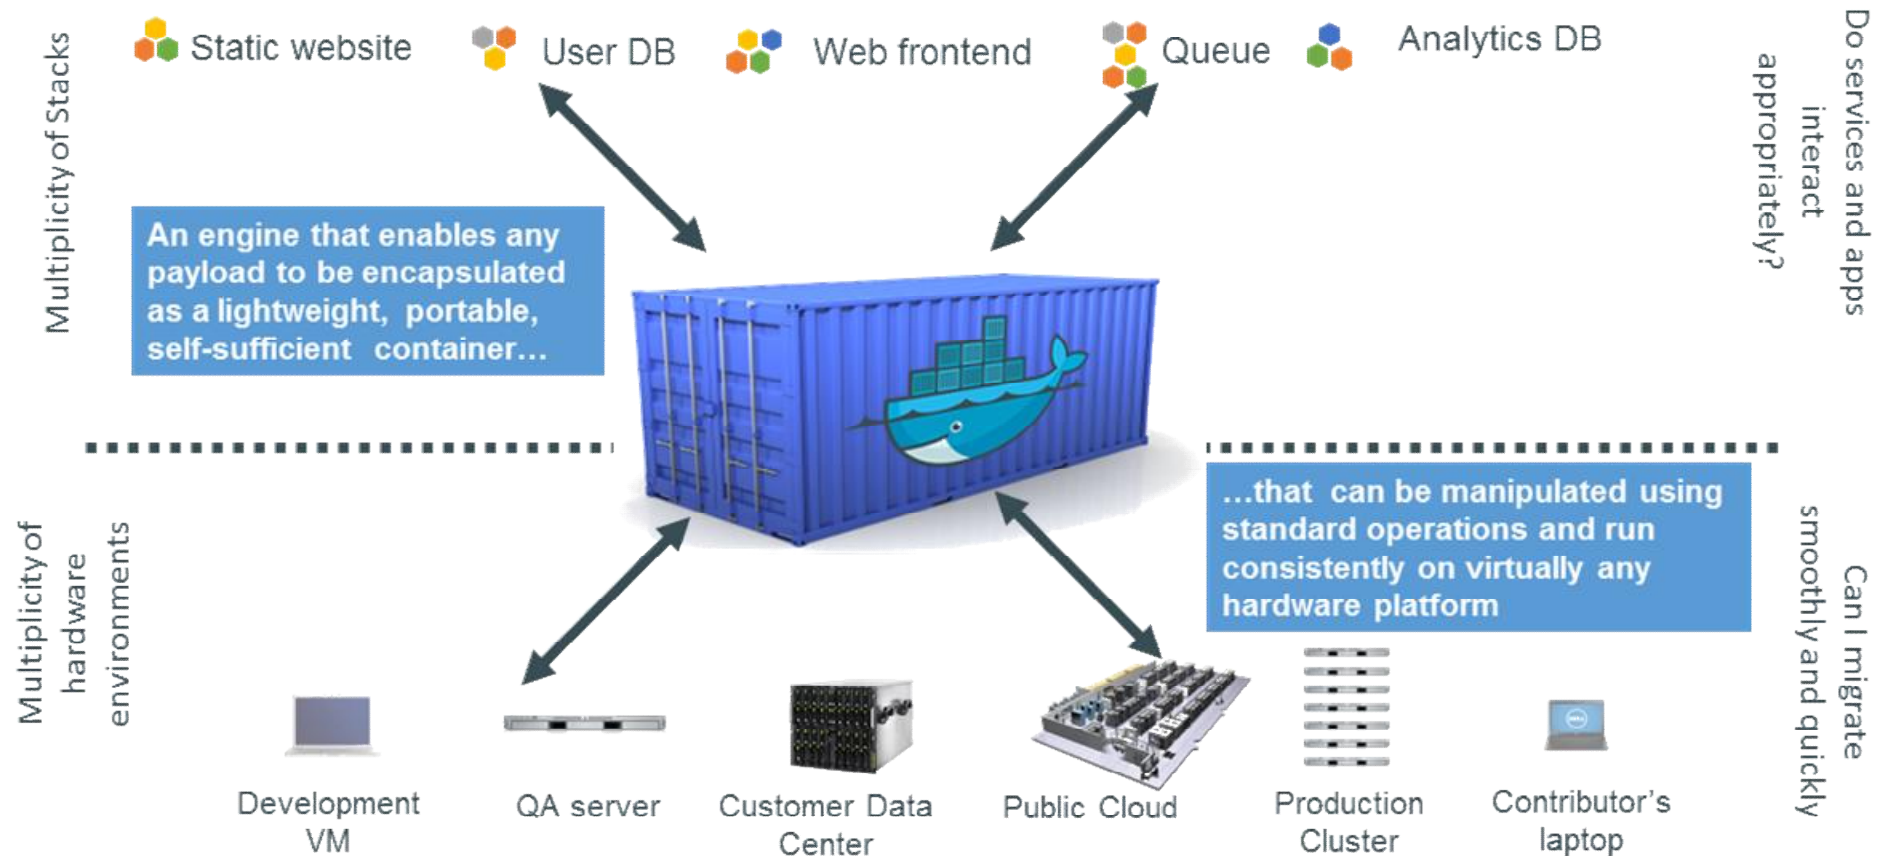
Also a Matrix from Hell

	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
							

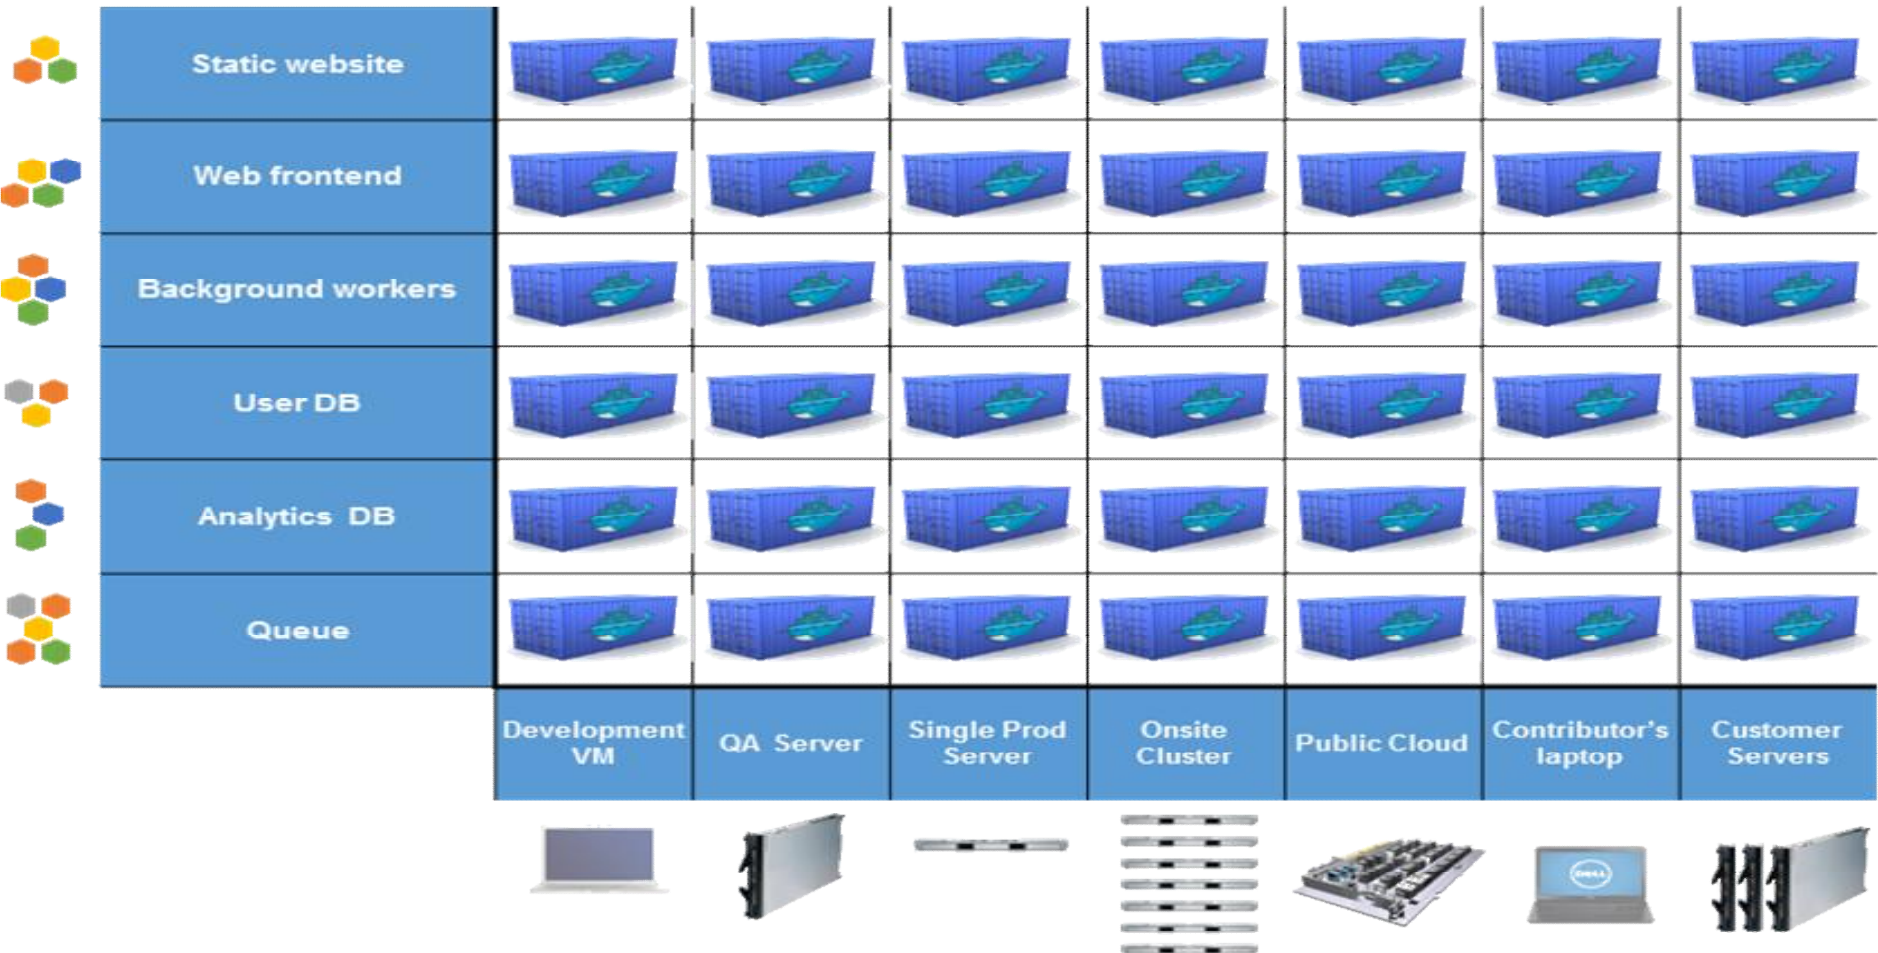
Solution: Intermodal Shipping Container



Docker is a Container System for Code

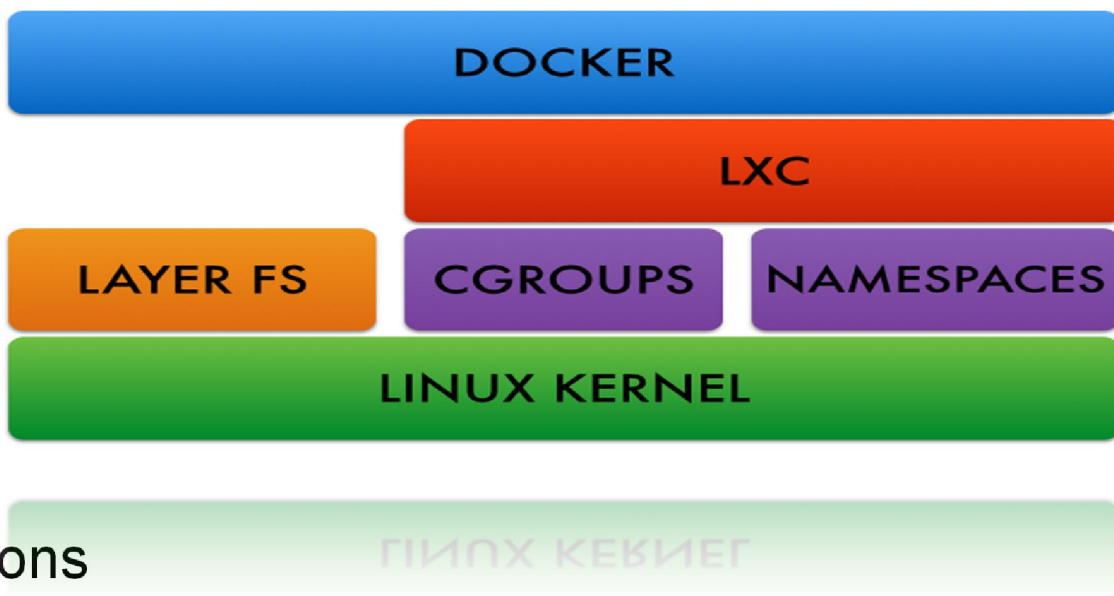


Docker Eliminates the Matrix from Hell



Docker Architecture

- Docker Engine
 - CLI
 - Docker Daemon
 - Docker Registry
- Docker Hub
 - Cloud service
 - Share Applications
 - Automate workflows
 - Assemble apps from components
- Docker images
- Docker containers



Docker Container

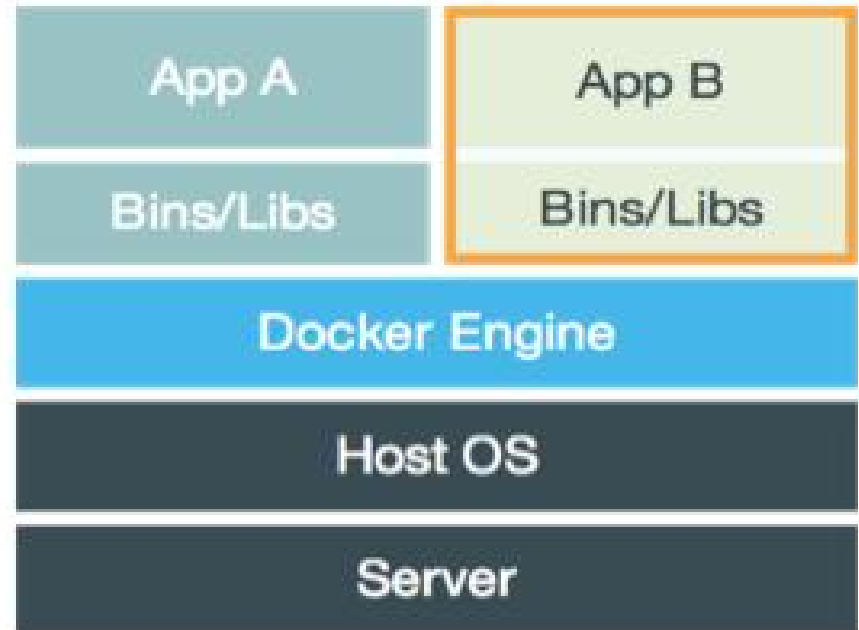
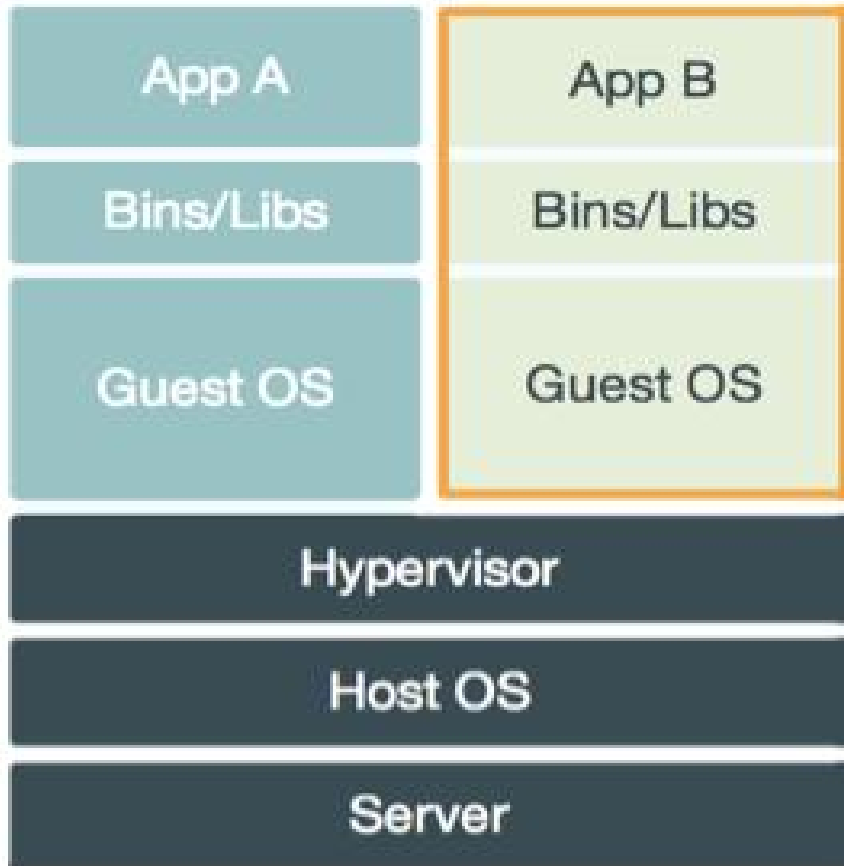
Units of software delivery (ship it!)

- Run everywhere
 - regardless of kernel version
 - regardless of host distro
 - (but container and host architecture must match*)
- Run anything
 - if it can run on the host, it can run in the container
 - i.e., if it can run on a Linux kernel, it can run

Docker Images

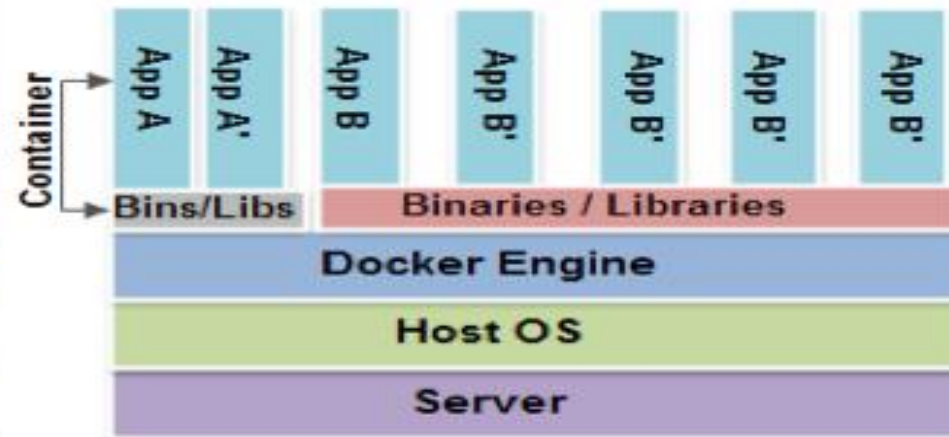
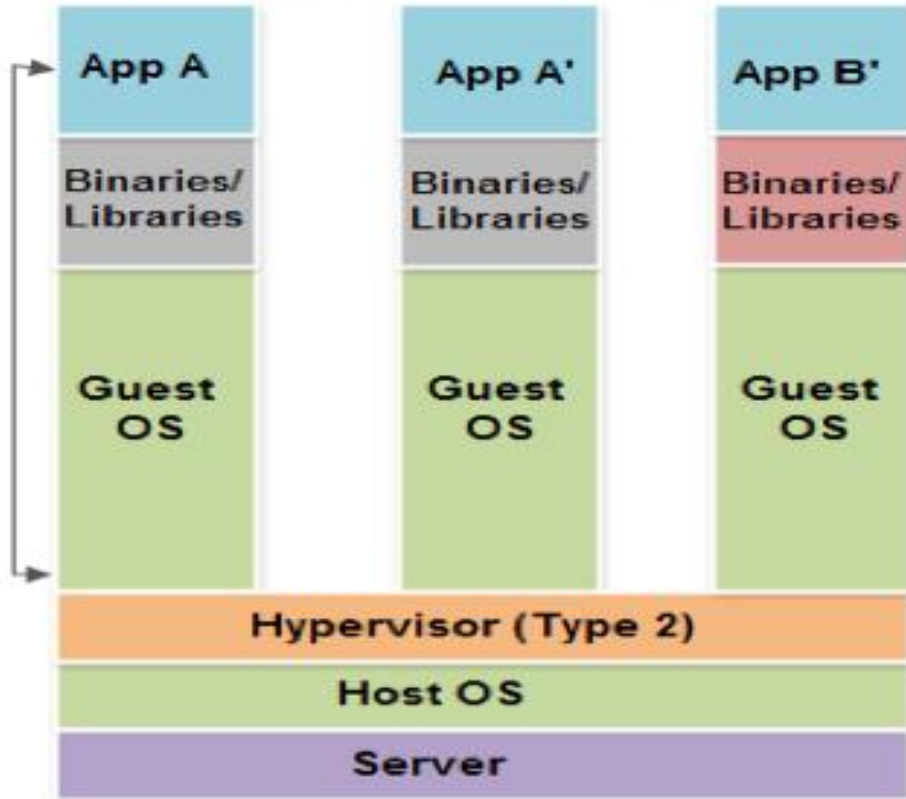
- NOT A VHD
- NOT A FILESYSTEM
- a read-only Layer
- do not have state
- Basically a tar file
- Has a hierarchy
 - Arbitrary depth
- Fits into the Docker Registry

Virtual Machine Versus Container



Virtual Machine Versus Container

Containers vs Virtual Machines



Docker Container Lifecycle

- The Life of a Container
 - Conception
 - **BUILD** an Image from a Dockerfile
 - Birth
 - **RUN** (create+start) a container
 - Reproduction
 - **COMMIT** (persist) a container to a new image
 - **RUN** a new container from an image
 - Sleep
 - **KILL** a running container
 - Wake
 - **START** a stopped container
 - Death
 - **RM** (delete) a stopped container
- Extinction
 - **RMI** a container image (delete image)